

# Introduction to **Information Retrieval**

Crawling and Duplicates

# This lecture

---


- Web Crawling
- (Near) duplicate detection

# Basic crawler operation

---

- Begin with known “seed” URLs
- Fetch and parse them
  - Extract URLs they point to
  - Place the extracted URLs on a queue
- Fetch each URL on the queue and repeat
- Breadth First crawling

# Crawling picture



# Simple picture – complications

---

- Web crawling isn't feasible with one machine
  - All of the above steps are usually distributed
- **Malicious pages**
  - Spam pages
  - Spider traps
- Even non-malicious pages pose challenges
  - Latency/bandwidth to remote servers vary
  - Webmasters' stipulations
    - How "deep" should you crawl a site's URL hierarchy?
  - Site mirrors and duplicate pages
- **Politeness – don't hit a server too often**

# What any crawler *must* do

---

- Be Polite: Respect implicit and explicit politeness considerations
- Be Robust: Be immune to spider traps and other malicious behavior from web servers

# Explicit and implicit politeness

---

- Explicit politeness: specifications from webmasters on what portions of a site can be crawled
  - robots.txt (see next slide)
- Implicit politeness: even with no specification, avoid hitting any site too often

# Robots.txt

---

- Protocol for giving spiders (“robots”) limited access to a website, originally from 1994
- Website announces its request on what can(not) be crawled
  - For a server, create a file `/robots.txt`
  - This file specifies access restrictions
- Details: [www.robotstxt.org/robotstxt.html](http://www.robotstxt.org/robotstxt.html)



# What any crawler *should* do

---

- Be capable of distributed operation: designed to run on multiple distributed machines
- Be scalable: designed to increase the crawl rate by adding more machines
- Performance/efficiency: permit full use of available processing and network resources

# What any crawler *should* do

---

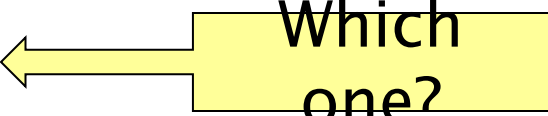
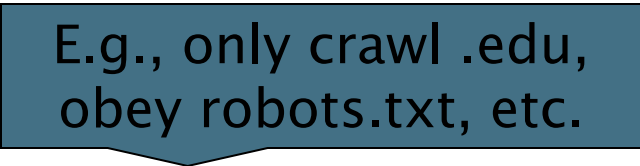
- Fetch pages of “higher quality” first
- Continuous operation: Continue fetching fresh copies of a previously fetched page
- Extensible: Adapt to new data formats, protocols

# URL frontier

---

- URLs that have been discovered, but are yet to be crawled
- Can include multiple pages from the same host
- **Must avoid trying to fetch them all at the same time**
- Must try to keep all crawling threads busy

# Processing steps in crawling

- Pick a URL from the frontier 
- Fetch the document at the URL
- Parse the URL
  - Extract links from it to other docs (URLs)
- Check if URL has content already seen
  - If not, add to indexes
- For each extracted URL 
  - Ensure it passes certain URL filter tests
  - Check if it is already in the frontier (duplicate URL elimination)

# Parsing: URL normalization

---

- When a fetched document is parsed, some of the extracted links are *relative* URLs
- E.g., [http://en.wikipedia.org/wiki/Main\\_Page](http://en.wikipedia.org/wiki/Main_Page) has a relative link to `/wiki/Wikipedia:General_disclaimer` which is the same as the absolute URL [http://en.wikipedia.org/wiki/Wikipedia:General\\_disclaimer](http://en.wikipedia.org/wiki/Wikipedia:General_disclaimer)
- During parsing, must normalize (expand) such relative URLs

# Content seen?

---

- Duplication is widespread on the web
- If the page just fetched is already in the index, do not further process it
- This is verified using document fingerprints or shingles
  - Second part of this lecture

# Distributing the crawler

---

- Run multiple crawl threads, under different processes – potentially at different nodes
  - May be geographically distributed nodes
- Partition hosts being crawled into nodes

## URL frontier: two main considerations

---

- Politeness: do not hit a web server too frequently
- Freshness: crawl some pages more often than others
  - E.g., pages (such as News sites) whose content changes often

These goals may conflict with each other.

(E.g., simple priority queue fails – many links out of a page go to its own site, creating a burst of accesses to that site.)



# Politeness – challenges

---

- Even if we restrict only one thread to fetch from a host, can hit it repeatedly
- Common heuristic: insert **time gap between successive requests to a host** that is  $\gg$  time for most recent fetch from that host

# Introduction to **Information Retrieval**

Near duplicate  
document detection

# Duplicate documents

---

- The web is full of duplicated content
- Strict duplicate detection = exact match
  - Not as common
- But many, many cases of **near duplicates**
  - E.g., Last modified date the only difference between two copies of a page

# Duplicate/Near-Duplicate Detection

---

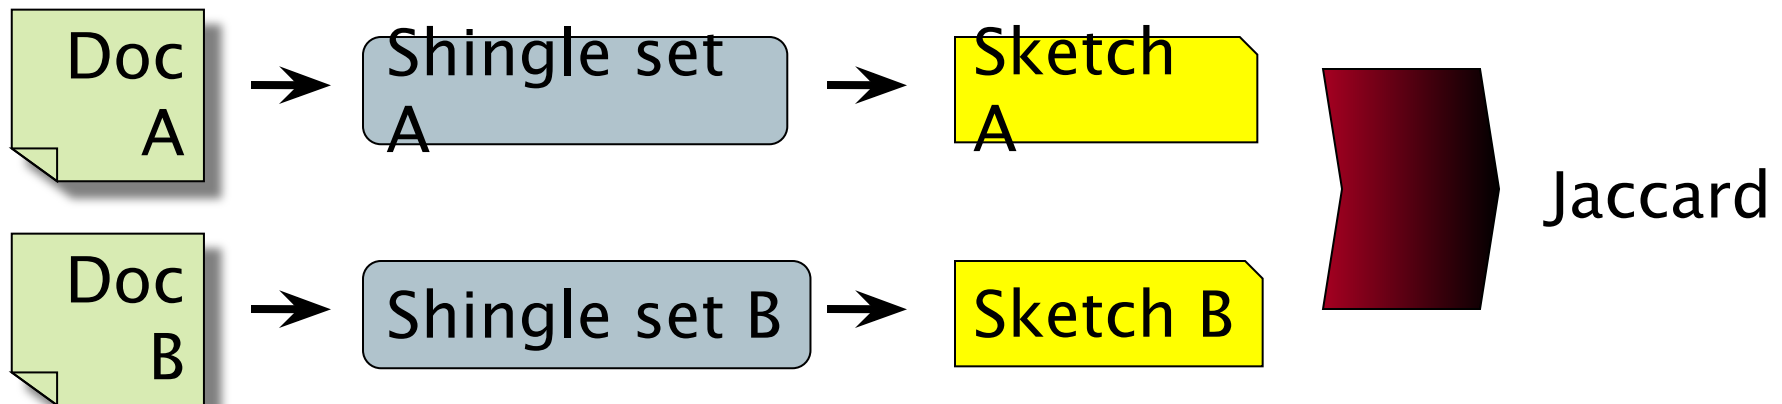
- *Duplication*: Exact match can be detected with fingerprints
  
- *Near-Duplication*: Approximate match
  - Overview
    - Compute syntactic similarity with an edit-distance measure
    - Use similarity threshold to detect near-duplicates, e.g.,  
Similarity > 80% => Documents are “near duplicates”

# Computing Similarity

- Features:
  - Segments of a document (natural or artificial breakpoints)
  - Shingles (Word N-Grams)
  - ***a rose is a rose is a rose*** → 4-grams are
    - a\_rose\_is\_a
    - rose\_is\_a\_rose
    - is\_a\_rose\_is
- Similarity Measure between two docs (= sets of shingles)
  - Jaccard coefficient: (Size\_of\_Intersection / Size\_of\_Union)

# Shingles + Set Intersection

- Computing exact set intersection of shingles between all pairs of documents is expensive
- Approximate using a cleverly chosen subset of shingles from each (a *sketch*)
- Estimate ( $\text{size\_of\_intersection} / \text{size\_of\_union}$ ) based on a short sketch



# Sketch of a document

---

- Create a “sketch vector” (of size  $\sim 200$ ) for each document
  - Documents that share  $\geq t$  (say 80%) corresponding vector elements are deemed **near duplicates**
  - For doc  $D$ ,  $\text{sketch}_D[i]$  is as follows:
    - Let  $f$  map all shingles in the universe to  $1..2^m$  (e.g.,  $f$  = fingerprinting)
    - Let  $\pi_i$  be a *random permutation* on  $1..2^m$
    - Pick  $\text{MIN} \{ \pi_i(f(s)) \}$  over all shingles  $s$  in  $D$

See details in book

# Final notes

---

- Shingling is a *randomized algorithm*
  - It will give us the right (wrong) answer with some probability on *any input*
- We've described how to detect near duplication in a pair of documents
- In “real life” we'll have to concurrently look at many pairs
  - See text book for details